

Scheduled Fetch

For this exercise, we're going to be applying what we've learned from the previous lessons in a simple DAG that fetches data from wikipedia.

The API we use gets the number of view of a wikipedia page. For this exercise we will get the number of wikipedia views of a particular page.

Pageviews API

The documentation for wikimedia api can be found here:

https://wikimedia.org/api/rest_v1/

The form of the request of pageview per articles is like this:

GET: [https://wikimedia.org/api/rest_v1/metrics/pageviews/per-article/en.wikipedia/all-access/all-agents/\[ARTICLE\]/daily/\[START\]/\[END\]](https://wikimedia.org/api/rest_v1/metrics/pageviews/per-article/en.wikipedia/all-access/all-agents/[ARTICLE]/daily/[START]/[END])

And the response is of the form:

```
{
  "items": [
    {
      "project": "en.wikipedia",
      "article": "Bitcoin",
      "granularity": "daily",
      "timestamp": "2015100100",
      "access": "all-access",
      "agent": "all-agents",
      "views": 7942
    },
    {
      "project": "en.wikipedia",
      "article": "Bitcoin",
      "granularity": "daily",
      "timestamp": "2015100200",
      "access": "all-access",
      "agent": "all-agents",
      "views": 7540
    }
  ]
}
```

The granularity for the pageview aggregation is "daily" pageviews.

We'll only be getting the pageviews per day, so the and date parameters must be the same day (`execution_date`).

Installing Libraries

To install any library for an airflow worker in `astro-cli`, you can put the library in the `requirements.txt` file in the root project directory.

`requirement.txt`

```
pandas
```

You can put other libraries you need here (ex. `scikit-learn`)

Then restart your `astro` project.

Dag Schedule

We'll schedule the DAG to start from 5 days ago and execute daily:

```
with DAG(  
    dag_id="5-scheduled-fetch",  
    start_date=airflow.utils.dates.days_ago(5),  
    schedule_interval="@daily"  
) as dag:
```

Fetch bitcoin views

We'll use `curl` on the `BashOperator` to make a GET request on the wikipedia pageviews endpoint.

```
fetch_bitcoin_views = BashOperator(  
    task_id="fetch_bitcoin_views",  
    bash_command="curl -o /tmp/{{ ds_nodash }}.json -L  
'https://wikimedia.org/api/rest_v1/metrics/pageviews/per-article/en.wikipedia/all-  
access/all-agents/Bitcoin/daily/{{ ds_nodash }}00/{{ ds_nodash }}00'"  
)
```

Notice how we use the `ds_nodash` (YYYYMMDD) template variable to both:

a.) Save the response as a json file. b.) API parameters, since wikimedia api needs the start and end date in this format.

Convert to CSV

We'll use the `pandas` library to load the json array of views into a dataframe, then save only the timestamp and views to the csv file.

```
def _save_to_csv(ds_nodash):  
    with open(f"/tmp/{ds_nodash}.json") as f:  
        views_json = json.load(f)  
        df = pd.DataFrame(views_json['items'])  
        df[['timestamp', 'views']].to_csv(f"/tmp/{ds_nodash}.csv", index=False)  
  
    ...  
  
    save_to_csv = PythonOperator(  
        task_id="save_to_csv",  
        python_callable=_save_to_csv,  
        provide_context=True  
    )
```

Challenge Exercise

Let's say that we want to correlate the views of the Bitcoin wikipedia page with it's price for that day.

You can use the [coindesk](https://www.coindesk.com/coindesk-api) api found here to get the closing price of bitcoin for a day:

<https://www.coindesk.com/coindesk-api>

Specifically, use this endpoint and modify the start and end date:

<https://api.coindesk.com/v1/bpi/historical/close.json?start=2013-09-01&end=2013-09-05>

Conver the bitcoin price for that day into CSV as well.

Finally, join both csv files according to their day:

End result: 1 CSV file with both the wikipedia views and the bitcoin price for that day.

Note:

You don't need to use [pandas](#) if you don't want to. As long as you have the final csv file of views and price.